

ASAPP Hackathon

Team Speed 'Speed Airbot'

23

October

2025

Team Members:
Kathir Kaarthik S
Keertanapriya D
Sarmitha VD
Srinidhi CR

Agenda

- | | | |
|----|---|----------|
| 01 | Problem Landscape | Slide 05 |
| 02 | Opportunity & Vision | Slide 08 |
| 03 | Technical Roadmap : From Vision to Value | Slide 09 |
| 04 | Target: Use Case Identification | Slide 10 |
| 05 | Setup: System Architecture | Slide 11 |
| 06 | Test: Conversation Flow Simulation | Slide 12 |

Agenda

07

Supervise: Monitoring & Quality Control

Slide 13

08

Scalability Beyond Airlines

Slide 14

Problem landscape – current scenario



**The Airline Support
Challenge : High Volume
and Low Efficiency in a
\$4.8 Billion market**

- Airlines handle **millions** of customer interactions daily across chat, calls, and emails.
- **Up to 65–70%** of these are **repetitive, rule-based requests** (e.g., flight status, refunds, cancellations).
- Despite existing self-service tools, **>60% of passengers** still prefer talking to human agents, citing lack of clarity and robotic responses.
- This creates **high operational costs, agent fatigue, and longer wait times.**

The pain points



✗ Inefficient manual resolution

Repetitive queries handled by agents



✗ Fragmented user experience

Bots fail to retain context or sound human



✗ Poor scalability

Customer spikes (flight delays, weather events) overwhelm systems



The Need

With global air travel back to pre-pandemic levels and digital-first engagement becoming standard, airlines face pressure to deliver instant, intelligent, and human-like support. Our solution bridges this gap with AI-native, policy-compliant conversational systems that blend automation and empathy.

To bridge this gap, our solution leverages ASAPP's data-driven roadmap to engineer an airline-grade conversational agent - one that's as intuitive as a human, yet as efficient as AI.

Vision & Opportunity



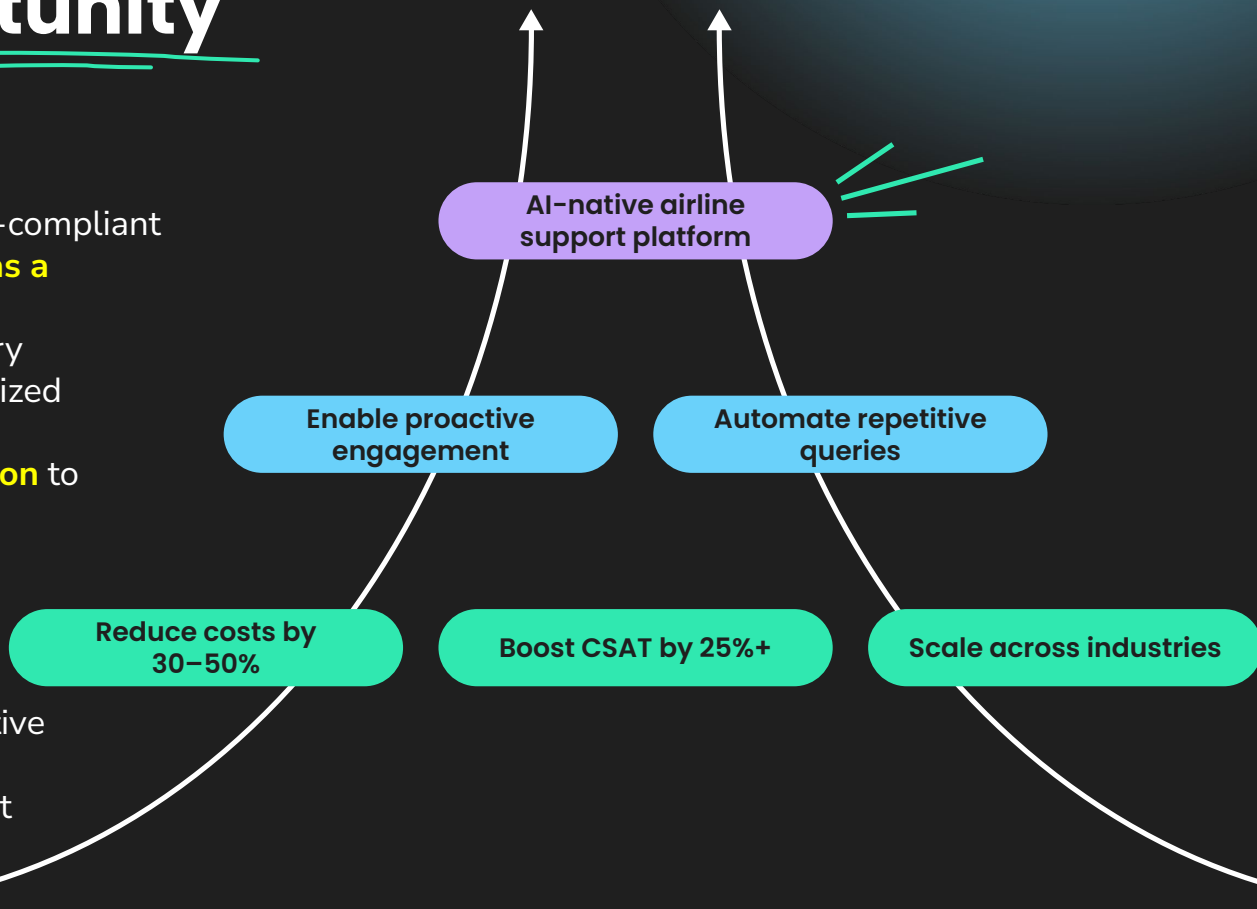
Our Vision

- Create a context-aware, policy-compliant AI agent that feels as **natural as a human**.
- Shift airlines from reactive query handling to proactive, personalized engagement.
- **Blend empathy with automation** to elevate customer experience.



Opportunity

- Airline contact-center market: \$4.8 B → \$12.6 B by 2033.
- **70 %** of interactions are repetitive and fit for automation.
- **Huge gap** in multi-turn, context retaining AI support solutions.



Technical Roadmap : From Vision to Value

Using ASAPP's 'Vision to Value' roadmap as our foundation, we've structured our solution to mirror its data-driven, scalable, and customer-centric approach

1. **Target** : Identify automation-heavy airline use cases
2. **Setup** : Integrate easily into airline systems
3. **Test** : Scale & simulate conversations safely
4. **Supervise** : Ensure compliance & consistency
5. **Run & Optimize** : Continuous learning and refinement

Target : Intelligent Use Case Prioritization

- **Booking Changes / Cancellations** → structured data retrieval + transaction.
- **Refund & Policy Queries** → repetitive and rule-based.
- **Baggage Tracking / Complaints** → moderate complexity with high query volume.
- **Loyalty Program Enquiries** → low urgency, consistent data flow.



Setup : System Design and Integration

- **Agents (Core Intelligence):**

- **Transactional Agent :** Executes API-based tasks (booking, refund, status).
- **Informational Agent :** Handles FAQ, policy lookups, and document retrieval via a vector database.
- Both use Relevance Scoring to pick the best response path.

- **Policy Module:**

- Ensures airline-specific tone, compliance, and wording.
- Filters off-topic or unsafe responses.

- **LLM Interface:**

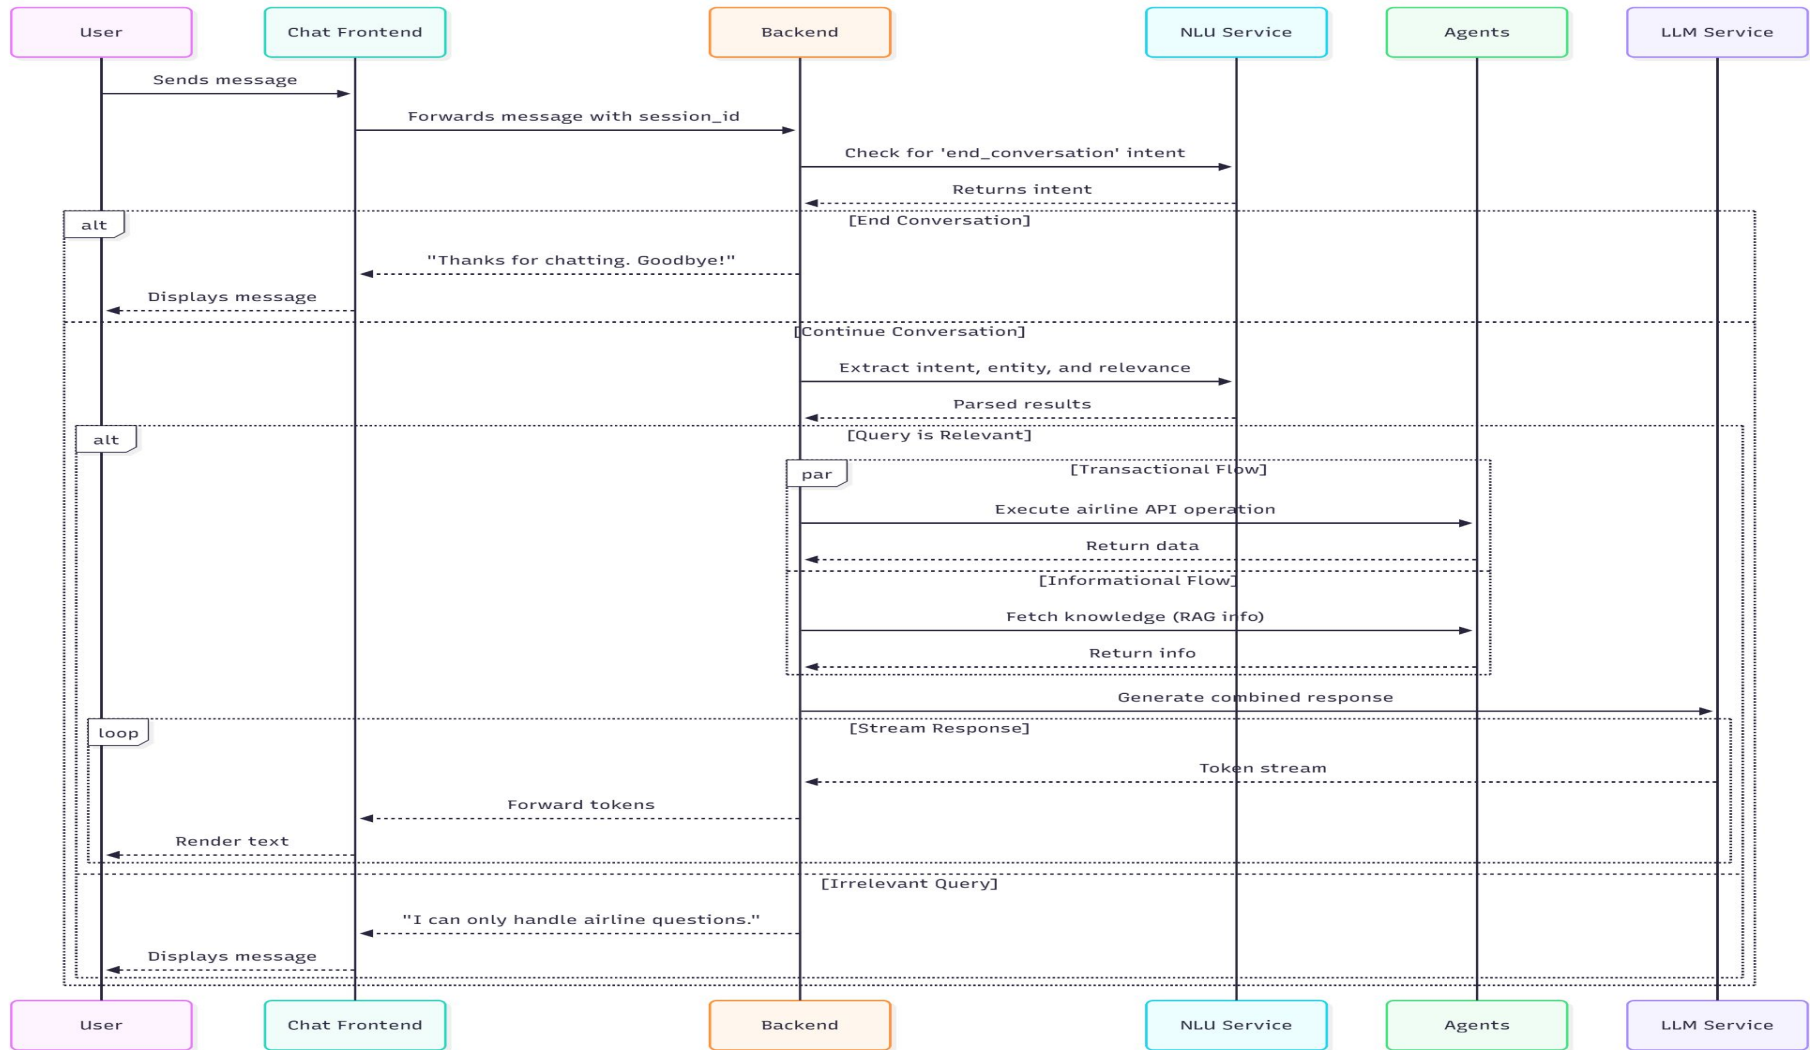
- Generates natural, human-like replies within airline domain boundaries.
- Integrated via secure API with response.

- **Frontend Layer:**

- Customer chat interface (web).
- Handles message input and session start.

- **Backend (Core Engine):**

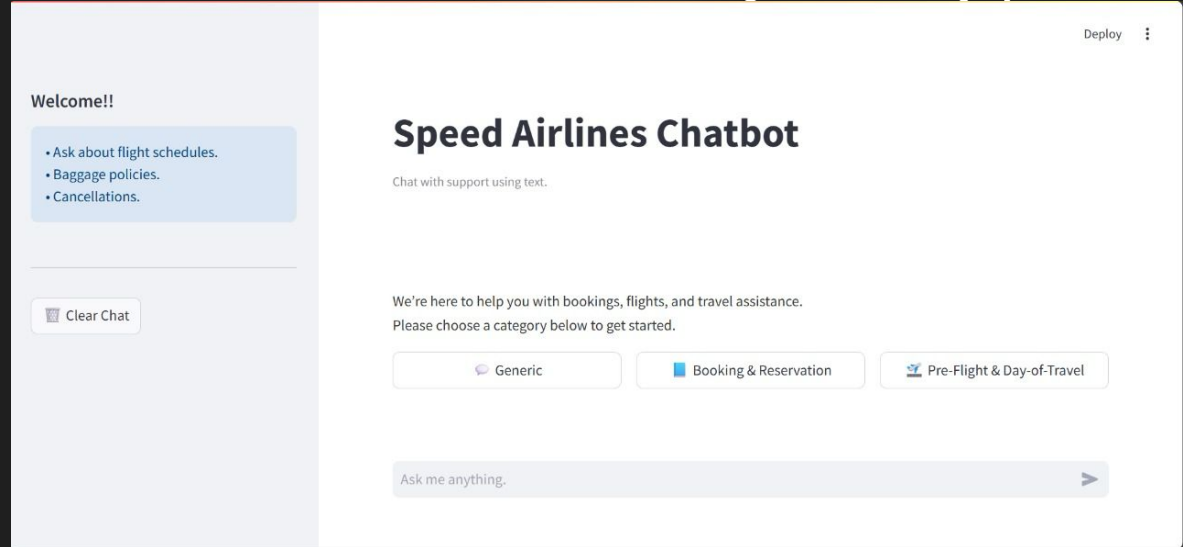
- **Session Manager :** tracks user context, multi-turn conversation state.
- **Redis Cache :** enables fast retrieval of previous turns and context memory.
- **Intent Classifier + Context Manager :** processes incoming messages, predicts intent, and maintains continuity across queries.



Test : Validation, Simulation & Quality Control

Key Testing Dimensions :


1. Intent Accuracy
2. Context Retention
3. Policy Adherence
4. Response Naturalness
5. Latency



Supervise : Monitoring & Quality Control

Goals we aim to achieve :

- Maintain **consistent response** quality across thousands of sessions.
- Detect and **flag policy violations** or off-topic responses in real time.
- Provide human oversight for **continuous** improvement and trust.








Scalability Beyond Airlines



Our modular architecture enables rapid adaptation across industries, transforming our airline support bot into a cross-sector conversational AI platform.

- Built with **domain-agnostic core modules** (session manager, NLU, policy engine).
- Industry-specific logic added via **plug-in intent models** and policy templates.
- Vector-based knowledge retrieval allows **easy retraining** with new datasets.
- API-first design supports **integration with any enterprise CRM or database.**

Industry	Example Use Case
Airlines	 Flight rebooking, baggage queries, refund policy assistance
Telecom	 Plan upgrades, outage reports, SIM activations
Banking	 Card blocking, transaction queries, KYC assistance
Retail / e-commerce	 Order tracking, returns, product availability
Logistics	 Shipment tracking, delay updates, delivery rescheduling

A white, curved, decorative line starts from the top left corner and extends diagonally towards the center of the slide.

Thank you!

[github](#)